

1 Overview

1.1 Location `$<APPSDKSamplesInstallPath>\samples\opencl\cpp_cl\`

1.2 How to Run See the *Getting Started* guide for how to build samples. You first must compile the sample.

Use the command line to change to the directory where the executable is located. The pre-compiled sample executable is at `$<APPSDKSamplesInstallPath>\samples\opencl\bin\x86\` for 32-bit builds, and `$<APPSDKSamplesInstallPath>\samples\opencl\bin\x86_64\` for 64-bit builds.

Type the following command(s).

1. `ConcurrentKernels`
This command runs the program with the default options.
2. `ConcurrentKernels -h`
This command prints the help file.

1.3 Command Line Options Table lists, and briefly describes, the command line options.
Command Line Options:

Short Form	Long Form	Description
-h	--help	Shows all command options and their respective meanings.
	--device	Devices on which the program is to be run. Acceptable values are <code>cpu</code> or <code>gpu</code> .
-q	--quiet	Quiet mode. Suppresses all text output.
-e	--verify	Verify results against reference implementation.
-t	--timing	Print timing-related statistics.
	--dump	Dump binary image for all devices.
	--load	Load binary image and execute on device.
-p	--platformId	Select platformId to be used (0 to N-1, where N is the number of available platforms).
-v	--version	AMD APP SDK version string.
-d	--deviceId	Select deviceId to be used (0 to N-1, where N is the number of available devices).
-i	--iterations	Number of iterations.
-x	--size	Size of the input buffer per kernel (in Bytes).

Short Form	Long Form	Description
-w	--workgroups	Number of workgroups per kernel execution.
-l	--localsize	Number of work items (2^N) per workgroup.
-k	--kernels	Number of kernels for concurrent executions.

2 Introduction

This sample shows how to execute more than one kernel concurrently on a GPU. Currently, concurrent kernel execution is supported on AMD's GCN architecture GPUs.

To run multiple kernels concurrently, the kernels must be submitted to different command queues. There must be no data or buffer dependency between the kernels. The GPU must possess enough compute resources to execute the kernels simultaneously.

The sample includes a sequential implementation in which kernel executions happen sequentially as only one command queue is used. Not all compute resources are fully utilized in this case. The concurrent version shows an improvement in the resource utilization of the GPU by way of executing kernels concurrently using separate command queues for each kernel.

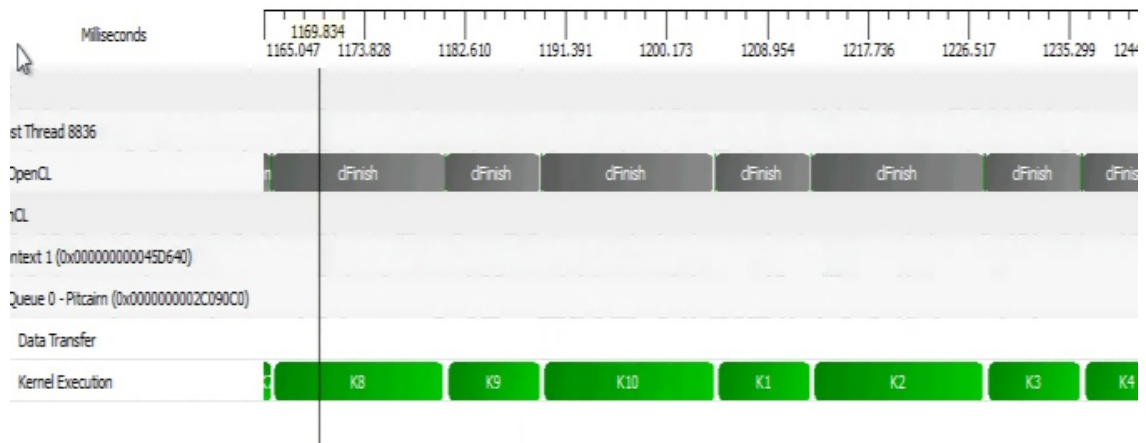
3 Implementation

The sample implements two versions of kernel executions:

1. Sequential
2. Concurrent


3.1 Sequential execution implementation

As the name says, the kernels run sequentially one after the other. The kernels are queued in a single queue and are executed sequentially. The execution of a kernel starts after the previous kernel completes execution. The following CodeXL snapshot illustrates the sequential operation:



3.2 Concurrent execution implementation

In this version, the kernels are executed concurrently by using separate command queues for each kernel. The following CodeXL snapshot illustrates the concurrent operation:



As seen, the kernels are queued in separate queues (1, 2, 3, and 4) and execute concurrently.

Contact

Advanced Micro Devices, Inc.
One AMD Place
P.O. Box 3453
Sunnyvale, CA, 94088-3453
Phone: +1.408.749.4000

For AMD Accelerated Parallel Processing:
URL: developer.amd.com/appsdk
Developing: developer.amd.com/
Support: developer.amd.com/appsdksupport
Forum: developer.amd.com/openclforum

AMD

The contents of this document are provided in connection with Advanced Micro Devices, Inc. ("AMD") products. AMD makes no representations or warranties with respect to the accuracy or completeness of the contents of this publication and reserves the right to make changes to specifications and product descriptions at any time without notice. The information contained herein may be of a preliminary or advance nature and is subject to change without notice. No license, whether express, implied, arising by estoppel or otherwise, to any intellectual property rights is granted by this publication. Except as set forth in AMD's Standard Terms and Conditions of Sale, AMD assumes no liability whatsoever, and disclaims any express or implied warranty, relating to its products including, but not limited to, the implied warranty of merchantability, fitness for a particular purpose, or infringement of any intellectual property right.

AMD's products are not designed, intended, authorized or warranted for use as components in systems intended for surgical implant into the body, or in other applications intended to support or sustain life, or in any other application in which the failure of AMD's product could create a situation where personal injury, death, or severe property or environmental damage may occur. AMD reserves the right to discontinue or make changes to its products at any time without notice.

Copyright and Trademarks
© 2013 Advanced Micro Devices, Inc. All rights reserved. AMD, the AMD Arrow logo, ATI, the ATI logo, Radeon, FireStream, and combinations thereof are trademarks of Advanced Micro Devices, Inc. OpenCL and the OpenCL logo are trademarks of Apple Inc. used by permission by Khronos. Other names are for informational purposes only and may be trademarks of their respective owners.

3 of 3

Concurrent Kernels