

1 Overview

1.1 Location `$<APPSDKSamplesInstallPath>\samples\opencl\cl\`

1.2 How to Run See the *Getting Started* guide for how to build samples. You first must compile the sample.

Use the command line to change to the directory where the executable is located. The pre-compiled sample executable is at `$<APPSDKSamplesInstallPath>\samples\opencl\bin\x86\` for 32-bit builds, and `$<APPSDKSamplesInstallPath>\samples\opencl\bin\x86_64\` for 64-bit builds.

Type the following command(s).

1. `Histogram`
Runs with default options; `x = 1024`, `y = 1024`.
2. `Histogram -h`
This prints the help file.

1.3 Command Line Options Table 1 lists, and briefly describes, the command line options.

Table 1 Command Line Options

| Short Form | Long Form | Description |
|------------|--------------|---------------------------------------------------------------------------------------------------------|
| -h | --help | Shows all command options and their respective meaning. |
| | --device | Devices on which the program is to be run. Acceptable values are <code>cpu</code> or <code>gpu</code> . |
| -q | --quiet | Quiet mode. Suppresses all text output. |
| -e | --verify | Verify results against reference implementation. |
| -t | --timing | Print timing. |
| | --dump | Dump binary image for all devices. |
| | --load | Load binary image and execute on device. |
| | --flags | Specify compiler flags to build the kernel. |
| -p | --platformId | Select platformId to be used (0 to N-1, where N is the number of available platforms). |
| -d | --deviceId | Select deviceId to be used (0 to N-1, where N is the number of available devices). |
| -v | --version | AMD APP SDK version string. |
| -x | --width | Width of the input array. |
| -y | --height | Height of the input array. |

| Short Form | Long Form | Description |
|------------|--------------|-----------------------------------------------------------------------------------------------------------|
| -i | --iterations | Number of iterations for kernel execution. |
| | --scalar | Run the scalar version of the kernel (Note that the --scalar and -vector options are mutually exclusive.) |
| | --vector | Run the vector version of the kernel (Note that the --scalar and -vector options are mutually exclusive.) |

2 Implementation Details

Each work-item calculates the histogram bin from BIN_SIZE elements. The block histogram bin are calculated from the histogram bins of the work-items of each block. The final histogram bin is calculated from the block histogram bins on the host (CPU) side.

The implementation contains two kernels: a scalar kernel and a vector kernel. The scalar kernel works better on current generation GCN cards. The vector kernel works better on VLIW and previous generation cards. The user can select the kernel using the --scalar or --vector options; otherwise, preferred vector-width for the device is queried and used. If both --scalar and --vector are specified, both the options are ignored and the default vector-width will be used.

Contact

Advanced Micro Devices, Inc.
One AMD Place
P.O. Box 3453
Sunnyvale, CA, 94088-3453
Phone: +1.408.749.4000

For AMD Accelerated Parallel Processing:
URL: developer.amd.com/appsdk
Developing: developer.amd.com/
Support: developer.amd.com/appsdksupport
Forum: developer.amd.com/openclforum



The contents of this document are provided in connection with Advanced Micro Devices, Inc. ("AMD") products. AMD makes no representations or warranties with respect to the accuracy or completeness of the contents of this publication and reserves the right to make changes to specifications and product descriptions at any time without notice. The information contained herein may be of a preliminary or advance nature and is subject to change without notice. No license, whether express, implied, arising by estoppel or otherwise, to any intellectual property rights is granted by this publication. Except as set forth in AMD's Standard Terms and Conditions of Sale, AMD assumes no liability whatsoever, and disclaims any express or implied warranty, relating to its products including, but not limited to, the implied warranty of merchantability, fitness for a particular purpose, or infringement of any intellectual property right.

AMD's products are not designed, intended, authorized or warranted for use as components in systems intended for surgical implant into the body, or in other applications intended to support or sustain life, or in any other application in which the failure of AMD's product could create a situation where personal injury, death, or severe property or environmental damage may occur. AMD reserves the right to discontinue or make changes to its products at any time without notice.

Copyright and Trademarks

© 2013 Advanced Micro Devices, Inc. All rights reserved. AMD, the AMD Arrow logo, ATI, the ATI logo, Radeon, FireStream, and combinations thereof are trademarks of Advanced Micro Devices, Inc. OpenCL and the OpenCL logo are trademarks of Apple Inc. used by permission by Khronos. Other names are for informational purposes only and may be trademarks of their respective owners.